

MRP

천정민



1 I. Recap

2 II. Method 1 - Monte-Carlo simulation

3 III. Method 2 - Iterative solution

I. Recap

Motivation

- I drink a bottle of soda everyday. I drink either Coke or Pepsi everyday. When I choose what to drink for today, I only consider what I drank yesterday.
- Specifically,
 - Suppose I drank Coke yesterday, then the chance of drinking Coke again today is 0.7.
 - (What is the chance of drinking Pepsi today then?)
 - Suppose I drank Pepsi yesterday, then the chance of drinking Pepsi again today is 0.5.
 - (What is the chance of drinking Coke today then?)
- Given I drink coke today, what is likely my consumption for upcoming 10 days? (Pepsi is \$1 and Coke is \$1.5)

- A Markov chain is a stochastic process with the specification of
 - a state space S
 - a transition probability matrix \mathbf{P}
- A Markov reward process is a Markov chain with the specification of
 - a reward r_t with the reward function $R(s) = \mathbb{E}[r_t | S_t = s]$
 - a time horizon H , which is the duration we are interested in cumulative sum of rewards.
- The return G_t is the sum of remaining reward at time t .
 - $G_0 = r_0 + r_1 + \dots + r_9$
 - $G_1 = r_1 + \dots + r_9$
 - $G_2 = r_2 + \dots + r_9$
 - ...
 - $G_9 = r_9$
- A state-value function $V_t(s)$ is the expected return given state s at time t . That is,
 $V_t(s) = \mathbb{E}[G_t | S_t = s]$

II. Method 1 - Monte-Carlo simulation

MC simulation for estimating state-value function

- Formally, for a finite-horizon MRP, the following is MC simulation for estimating state-value function.

```
# MC evaluation for state-value function
# with state s, time 0, reward r, time-horizon H
1: episode_i <- 0
2: cum_sum_G_i <- 0
3: while episode_i < num_episode
4:   Generate an stochastic path starting from state s and time 0.
5:   Calculate return  $G_i$  <- sum of rewards from time 0 to time H-1.
6:   cum_sum_G_i <- cum_sum_G_i +  $G_i$ 
7:   episode_i <- episode_i + 1
8: State-value-fn  $V_t(s)$  <- cum_sum_G_i/num_episode
9: return  $V_t(s)$ 
```

- Remark that the full stochastic evolution, previously marked as MC_i is replaced by the term episode_i. Episode refers to a full single stochastic path from now on.

III. Method 2 - Iterative solution

Motivation

- Same as the previous section, our goal is still to estimate $V_0(c) = \mathbb{E}[G_0|S_t = c]$.
- Since $G_t = \sum_{i=t}^9 r_i$ has less number of terms when t is high number, we shall start from $t = 9$ and work backward, i.e. from $V_9(s)$, then $V_8(s)$, then $V_7(s)$, ...
- For $t = 9$,
 - From the general formula $V_t(s) = \mathbb{E}[G_t|S_t = s]$, it is easy to see that $V_9(s) = \mathbb{E}[G_9|S_9 = s] = \mathbb{E}[\sum_{i=9}^9 r_i|S_9 = s] = \mathbb{E}[r_9|S_9 = s] = R(s)$.
 - In other words,
 - $V_9(c) = \mathbb{E}[r_9|S_9 = c] = R(c) = 1.5$ and
 - $V_9(p) = \mathbb{E}[r_9|S_9 = p] = R(p) = 1.0$.
 - In general,

$$V_9(s) = R(s) + V_{10}(s), \quad (1)$$

where $V_{10}(s) = 0, \forall s$

- For $t = 8$,

- From the general formula $V_t(s) = \mathbb{E}[G_t | S_t = s]$, (watch below carefully)

$$\begin{aligned}
 V_8(s) &= \mathbb{E}[G_8 | S_8 = s] \\
 &= \mathbb{E}\left[\sum_{i=8}^9 r_i \mid S_8 = s\right] \\
 &= \mathbb{E}[r_8 + r_9 | S_8 = s] \\
 &= \mathbb{E}[r_8 | S_8 = s] + \mathbb{E}[r_9 | S_8 = s] \\
 &= R(s) + \mathbb{E}[r_9 | S_8 = s]
 \end{aligned} \tag{2}$$

- Here, let's consider $\mathbb{E}[r_9 | S_8 = c]$ first.

- This is expected spending on day-9 given that I drink coke on day-8. This value is conditioned on what I drink on day-9. If coke on day-9 with probability 0.7, $r_9 = 1.5$. If pepsi w/ prob. 0.3, $r_9 = 1.0$. This expectation is $1.35 (= 0.7 \cdot 1.5 + 0.3 \cdot 1.0)$.
- Formally, (using $\mathbb{E}(X|A) = \mathbb{E}(X|E_1, A)\mathbb{P}(E_1|A) + \mathbb{E}(X|E_2, A)\mathbb{P}(E_2|A)$)

$$\begin{aligned}
 &\mathbb{E}[r_9 | S_8 = c] \\
 = &\mathbb{E}[r_9 | S_9 = c, S_8 = c]\mathbb{P}(S_9 = c | S_8 = c) + \mathbb{E}[r_9 | S_9 = p, S_8 = c]\mathbb{P}(S_9 = p | S_8 = c) \\
 = &\mathbf{P}_{cc}\mathbb{E}[r_9 | S_9 = c] + \mathbf{P}_{cp}\mathbb{E}[r_9 | S_9 = p] \quad (\because \text{Markov property}) \\
 = &\mathbf{P}_{cc}\mathbb{E}[G_9 | S_9 = c] + \mathbf{P}_{cp}\mathbb{E}[G_9 | S_9 = p] = \mathbf{P}_{cc}V_9(c) + \mathbf{P}_{cp}V_9(p)
 \end{aligned}$$

- (Cont'd for $t = 8$)

- Now, let's consider $\mathbb{E}[r_9|S_8 = s]$ for the generalized state s . With the notation assuming a transition from this state s to the next state s' ,

$$\begin{aligned}\mathbb{E}[r_9|S_8 = s] &= \mathbf{P}_{sc}V_9(c) + \mathbf{P}_{sp}V_9(p) \\ &= \sum_{s' \in S} \mathbf{P}_{ss'}V_9(s')\end{aligned}\tag{3}$$

- We shall now summarize for $t = 8$,

$$\begin{aligned}V_8(s) &= \mathbb{E}[G_8|S_8 = s] = \mathbb{E}[r_8 + G_9|S_8 = s] \\ &= R(s) + \mathbb{E}[G_9|S_8 = s] \\ &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'}V_9(s')\end{aligned}\tag{4}$$

(expected return at time 8) = (reward at time 8) + (expected return at time 9)

- For $t = 7$,
 - From the general formula $V_t(s) = \mathbb{E}[G_t | S_t = s]$,

$$\begin{aligned} V_7(s) &= \mathbb{E}[G_7 | S_7 = s] \\ &= \mathbb{E}\left[\sum_{i=7}^9 r_i \mid S_7 = s\right] \\ &= \mathbb{E}[r_7 + r_8 + r_9 | S_7 = s] \\ &= \mathbb{E}[r_7 | S_7 = s] + \mathbb{E}[r_8 + r_9 | S_7 = s] \\ &= R(s) + \mathbb{E}[G_8 | S_7 = s] \end{aligned} \tag{5}$$

- You got the hint? From here, we want to use $V_8(s) = \mathbb{E}[G_8 | S_8 = s]$ to express this as a recursive formula for state-value function just like Eq. (4).

$$\begin{aligned} V_7(s) &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} \mathbb{E}[G_8 | S_7 = s, S_8 = s'] \\ &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_8(s') \end{aligned} \tag{6}$$

- So far,

$$\begin{aligned}V_{10}(s) &= 0 \\V_9(s) &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_{10}(s') \text{ from Eq. (1)} \\V_8(s) &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_9(s') \text{ from Eq. (4)} \\V_7(s) &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_8(s') \text{ from Eq. (6)} \\&\dots = \dots \\V_t(s) &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_{t+1}(s') \\&\dots = \dots \\V_0(s) &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_1(s')\end{aligned}$$

- Note that the array of equations can be solve from the top to the bottom.
- This iterative method is called as backward induction that works well with finite horizon problem.
- This iterative method (and its painful derivaion) is the most important mathematical essence of Markov decision process.

Implementation strategy

- Summary so far

$$\begin{aligned}V_{10}(s) &= 0 \\ V_t(s) &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_{t+1}(s') \quad (\text{for } t \in \{0, 1, \dots, 9\})\end{aligned}$$

- Strategy

- Column vector v_t for $V_t(s)$
- Column vector R for $R(s)$
- The term $\sum_{s' \in S} \mathbf{P}_{ss'} V_{t+1}(s')$ can be written as $\mathbf{P}v_{t+1}$.
- It follows

$$v_t = R + \mathbf{P}v_{t+1},$$

simply a system of linear equations!

```
import numpy as np

# Transition matrix P (R fills by column; equivalent Python Layout shown)
P = np.array([[0.7, 0.3],
              [0.5, 0.5]], dtype=float)

# Reward vector R
R = np.array([[1.5],
              [1.0]], dtype=float)

# Time-horizon
H = 10

#  $v_{t+1}$  initialization (terminal value is zero)
v_t1 = np.zeros((2, 1), dtype=float)

t = H - 1
while t >= 0:
    # Bellman recursion:  $v_t = R + P * v_{t+1}$ 
    v_t = R + P @ v_t1
    # step backward
    t -= 1
    # shift for next iteration
    v_t1 = v_t

#  $v_t$  now holds  $v_{\theta}(c)$ ,  $v_{\theta}(p)$  as a column vector
print(v_t)

## [[13.35937498]
##  [12.73437504]]
```

"0813-Jeongmin"